You are testing an application that uses EC2 instances to poll an SQS queue. At this stage of testing, you have verified that the EC2 instances can retrieve messages from the queue, but your coworkers are complaining of not being able to manually retrieve any messages from the queue from their on-premises workstations. What is the most likely source of this problem?"

- Your coworkers may not have permission to the SQS queue
- They are using short polling
- Manual polling of SQS queues is not possible
- SQS queues only accept traffic from within AWS

Explanation

When a snapshot is being taken against an EBS volume, the volume becomes unavailable and the instance no longer has the ability to communicate with the EBS volume until the snapshot is complete.

- False
- True

Explanation

Which is an operational process performed by AWS for data security?

- Background virus scans of EBS volumes and EBS snapshots
- Secure wiping of EBS data when an EBS volume is unmounted
- Replicating Data over multiple AWS Regions
- Decommissioning of storage devices using industry-standard practices

Explanation

Your AWS environment contains several reserved EC2 instances dedicated to a project that has just been cancelled. Your supervisor would like to recuperate cost for these reserved instances, but also does not want to lose the data just yet in case the project is revived next fiscal year. What can you do to minimize charges for these instances? Choose the 2 correct answers:

- Contact AWS and explain the situation
- Take snapshots of the EBS volumes and terminate the instances
- Stop the instances as soon as possible
- Sell the instances on the AWS Reserved Instance Marketplace

Explanation

To prevent in-flight tampering, all requests sent with API keys over REST or Query API should be sent over HTTPS connection.

- False
- True

Explanation

Which of the following relational database servers are available on Amazon RDS? Choose the 3 correct answers:

- Aurora
- MySQL
- DB2
- MariaDB

Explanation

You recently purchased and deployed four reserved EC2 instances in the US-East-1 region's Availability Zone 1 for a new project. Your supervisor just informed you that this project only requires two EC2 instances. Rather than selling the reserved instances, she asked you to terminate the extra instances and convert two of the on-demand instances already running in Availability Zone 1 to reserved instances. Can this be done?

- Yes, you can request this billing change through an AWS form.
- Yes, you can terminate the on-demand instances and re-deploy them as reserved instances.
- No, it cannot be done.
- Yes, you can terminate the reserved instances and AWS will automatically begin billing the two on-demand instances as reserved instances.

Explanation

Your AWS environment contains several on-demand EC2 instances dedicated to a project that has just been cancelled. Your supervisor does not want to incur charges for these on-demand instances, but also does not want to lose the data just yet because there is a chance the project may be revived in the next few days. What should you do to minimize charges for these instances in the meantime?

- Sell the instances on the AWS On-Demand Instance Marketplace. You can buy them back later if needed
- Stop the instances as soon as possible
- Contact AWS and explain the situation
- Terminate the instances as soon as possible

Explanation

If an instance, that belongs to an Elastic Load balancer, health check fails; what occurs to the instance that fails?

- The ELB will continue serving the traffic until the instance is removed
- The ELB will de-register the instance and stop sending traffic to the unhealthy instance
- The Auto Scaling group will launch a new instance
- ELB launches a new healthy instance

Explanation

Which of the following will occur when an EC2 instance in a VPC (Virtual Private Cloud) with an associated Elastic IP is stopped and started? Choose the 2 correct answers:

- The Elastic IP will be dissociated from the instance
- The ENI (Elastic Network Interface) is detatched
- All data on instance-store devices will be lost in instance store backed image
- The underlying host for the instance could be changed

Explanation

You design an application that checks for new items in an S3 bucket once per hour. If new items exist, a message is added to an SQS queue. You have several EC2 instances which retrieve messages from the SQS queue, parse the file, and send you an email containing the relevant information from the file. You upload one test file to the bucket, wait a couple hours and find that you have hundreds of emails from the application. What is the most likely cause for this volume of email?

- This is expected behavior when using long polling because SQS does not guarantee that there will not be duplicate messages processed.
- This is expected behavior when using short polling because SQS does not guarantee that there will not be duplicate messages processed.
- You can only have one EC2 instance polling the SQS queue at a time.
- Your application does not issue a delete command to the SQS queue after processing the message.

Explanation

An EC2 instance retrieves a message from an SQS queue, begins processing the message, then crashes. What happens to the message?

- To prevent data loss, it remains in the queue in a locked state until the EC2 instance comes back online.
- To prevent data corruption, if the EC2 instance does not respond to a customizable number of pings, the message is deleted.
- When the message visibility timeout expires, the message becomes available for processing by other EC2 instances.
- To prevent data corruption, when the message hide timeout expires, the message is duplicated, the original message is archived, and the duplicate message becomes available for processing by other EC2 instances.

Explanation

Which of the following is not a legitimate concern?

- Multi-AZ RDS deployments require at least two Availability Zones.

- Databases typically should not be deployed on a public subnet.
- US-East-1 is does not support Multi-AZ RDS deployments.
- Multi-AZ RDS deployments require a minimum of two subnets in a subnet group.

Explanation

Your company is moving their entire 20TB data warehouse to the cloud. With your current bandwidth it would take 2 months to transfer the data. Which service would allow you to quickly get your data into AWS?

- Amazon S3 Connector
- Amazon S3 MultiPart Upload
- Amazon Import/Export
- Amazon Direct Connect

Explanation

Your company has moved a legacy application from on-premise data centre to the cloud. The legacy application requires a static IP address being coded into the application which prevents you from deploying the application with high availability and fault tolerance using the ELB. Which steps would you take to apply high availability and fault tolerance to this application? Choose the 2 correct answers:

- Write a custom script that pings the health of the instance and if the instance stops responding, switches the elastic IP address to a standby instance
- Create an AMI of the instance and launch it using Auto Scaling which will deploy the instance again if it becomes unhealthy
- Do not migrate the application to the cloud until it can be converted to work with the ELB and Auto Scaling
- Ensure that the instance it's using has an elastic IP address assigned to it

Explanation

You manage a web application on EC2 instances. Your website occasionally experiences brief but large spikes in traffic that cause your EC2 instances' resources to become overwhelmed and the application to freeze up and lose recently submitted requests from end users. You use Auto Scaling to deploy additional resources to handle the load during spikes, but the new instances do not deploy fast enough to prevent the existing application servers from freezing. Assuming you cannot find a pattern to when these spikes occur, which of the following is likely to provide the most cost-effective solution to avoid losing recently submitted requests?

- Keep one extra EC2 instance always powered on in case a spike occurs.
- Pre-warm your Elastic Load Balancer.
- Use larger instances for your application.
- Use an SQS queue to decouple the application components.

Explanation

After deciding that due to strict contractual requirements that the latest AWS VPC that you deploy will need to incorporate AWS CloudHSM as an encryption solution. Where within your AWS infrastructure would be the best place to physically locate the HSM appliances and why?

- Locating HSM appliances completely isolated from your EC2 instances in another region increases secur
- Locating HSM appliances near your EC2 instances decreases network latency, which can improve application performance.
- Locating Locating HSM appliances in a private subnet increases security
- Locating HSM appliances nearest to the majority of potential customers decreases network latency, which can improve application performance.

Explanation

One of your instances is reporting an unhealthy system status check. However, this is not something you should have to monitor and repair on your own. How might you automate the repair of the system status check failure in an AWS environment?

- Create CloudWatch metrics that stop and then start the instance based off of status check alarms
- Write a script that queries the EC2 API for each instance status check
- Write a script that periodically shuts down and starts instances
- Implement a third party monitoring tool such as Nagios

Explanation

The AMI ID used in an Auto Scaling policy is configured in the _____.

- group policy
- launch configuration
- Auto Scaling group
- Auto Scaling Policy

Explanation

Which of the following is not a benefit of a decoupled architecture using EC2, Auto Scaling, and SQS?

- An application does not become unavailable due to the deletion of a single SQS queue
- Unprocessed SQS messages are not lost due to the failure of a single EC2 instance
- An application does not become unavailable due to the failure of a single EC2 instance
- B. Auto Scaling can spin up new EC2 instances to handle an increased number of items in an SQS queue

Explanation

Your AWS environment contains several reserved EC2 instances dedicated to a project that has just been cancelled. Your supervisor wants to stop incurring charges for these reserved instances immediately and recuperate as much of the reserved instance cost as possible. What can you do to avoid being charged for them? Choose the 2 correct answers:

- Stop the instances as soon as possible
- Terminate the instances as soon as possible
- Sell the reserved instances on the AWS Reserved Instance Marketplace
- Contact AWS and explain the situation

Explanation

What URL might you query on an EC2 instance in order to find the public AND private IP address of an instance?

- http://169.254.169.169/latest/meta-data/
- http://169.254.169.254/latest/meta-data/
- http://169.254.169.254/latest/user-data/
- http://169.254.169.169/latest/meta-data/

Explanation

In order to establish a successful site-to-site VPN connection from your on-premise network to the VPC (Virtual Private Cloud), which of the following needs to be configured inside of the VPC?

- Modify the main route table to route traffic through a NAT instance
- Need to launch a dedicated NAT instance in a public subnet
- Assign an Elastic IP address to the Virutal Private Gateway
- Assign a public IP address on the customer gateway for the on-premise network

Explanation

Your company is concerned with EBS volume backup on Amazon EC2 and wants to ensure they have proper backups and that the data is durable. What solution would you implement and why?

- Configure Amazon Storage Gateway with EBS volumes as the data source and store the backups on premise through the storage gateway
- Write a cronjob that uses the AWS CLI to take a snapshot of production EBS volumes. The data is durable because EBS snapshots are stored on the Amazon S3 standard storage class
- Use a lifecycle policy to backup EBS volumes stored on Amazon S3 for durability
- Write a cronjob on the server that compresses the data that needs to be backed up using gzip compression then use AWS CLI to copy the data into an S3 bucket for durability

Explanation

While implementing a disaster recovery strategy in another region, you are attempting to move the data from one EBS volume to another in a separate region. What is the best way to do this? Keep in mind this is not a live production replication copy.

- Configure rsync on your instance and sync the data across regions to another configured EC2 instance
- Take a snapshot of the EBS volume and copy it to the desired region
- Configure VPC peering to setup a direct link and copy a snapshot of the volume
- Copy the snapshot from S3 by enabling S3 region replication

Explanation

You are the System Administrator for your company's AWS account of approximately 200 IAM users. A new company policy has just been introduced that will change the access of 50 of the IAM users to have unlimited access to S3 buckets. How can you implement this effectively so that there is no need to apply the policy at the individual user level?

- Create a new role and add each user to the IAM role
- Create a policy and apply it to multiple users using a JSON script
- Use the IAM groups and add users as per their role to different groups and apply the policy to group
- Create an S3 bucket policy with unlimited access which includes each user's AWS account ID

Explanation

"Which of the following AWS services allow you access to the underlying operating system? Choose the 2 correct answers:"

- Amazon EMR
- Amazon Elasticbeanstalk
- Amazon S3
- Amazon RDS

Explanation

Amazon Auto Scaling is not meant to handle instant load spikes but is built to grow with a gradual increase in usage over a short time period.

- True
- False

Explanation

"After building an application that makes use of an Elastic Load balancer over port 80 you notice that your instances, even though they are healthy as part of the health check, are not serving traffic when you go to the ELB DNS cname. What might be the cause of this issue? Choose the 2 correct answers:"

- The EC2 instances are part of a public subnet
- The EC2 instances do not have port 80 open
- There is no attached internet gateway
- The ELB security group does not have port 80 open

Explanation

Your company wants to backup the onsite file server to AWS but does not want to serve the files from S3 to your office network when files need accessed. Which service and setup would you use to accomplish this task?

- Use Amazon Import/export
- Configure an EBS volume, open the security group to allow for NFS, and mount the volume on the local network
- Use Amazon Storage Gateway and gateway-cached volumes to store the data locally and asynchronously backup point-in-time snapshots to S3
- Use Amazon Storage Gateway and gateway-stored volumes to store the data locally and asynchronously backup point-in-time snapshots to S3

Explanation

To protect S3 data from both accidental deletion and accidental overwriting, you should:

- access S3 data using only signed URLs
- enable Multi-Factor Authentication (MFA) protected access
- enable S3 versioning on the bucket
- disable S3 delete using an IAM bucket policy

Explanation

Your web application front end consists of multiple EC2 instances behind an Elastic Load Balancer. You configured ELB to perform health checks on these EC2 instances. If an instance fails to pass health checks, which statement will be true?

- The instance gets terminated automatically by the ELB.
- The instance is replaced automatically by the ELB.
- The instance gets quarantined by the ELB for root cause analysis.
- The ELB stops sending traffic to the instance that failed its health check

Explanation

Your supervisor asks you to create a highly available website which serves static content from EC2 instances. Which of the following is not a requirement to accomplish this goal?

- An SQS queue
- An Auto Scaling group to recover from EC2 instance failures
- Multiple availability zones
- Multiple subnets

Explanation

Which of the following is true of an SQS message?

- SQS messages can live in the queue up to thirty days
- SQS messages must be less than 32 KB in size
- SQS messages are guaranteed to be delivered at least once
- SQS messages must be in JSON format

Explanation

You are building a system to distribute confidential training videos to employees. Using CloudFront, what method would be used to serve content that is stored in S3, but not publicly accessible from S3 directly?

- Create an Identity and Access Management (IAM) User for CloudFront and grant access to the objects in your S3 bucket to that IAM user
- Create a S3 bucket policy that lists the CloudFront distribution ID as the Principal and the target bucket as the Amazon Resource Name (ARN)
- Create an Origin Access Identify (OAI) for CloudFront and grant access to the objects in your S3 bucket to that OAI.
- Add the CloudFront account security group

Explanation

What is the difference between an availability zone and an edge location?

- Edge locations are used as control stations for AWS resources
- An availability zone is a grouping of AWS resources in a specific region; an edge location is a specific resource within the AWS region
- An availability zone is an Amazon Resource within an AWS region whereas an edge location will deliver cached content to the closest location to reduce latency
- None of the above

Explanation

Your supervisor asks you to create a highly available, decoupled web application. Which of the following does not help you accomplish this goal?

- An SQS queue to allow a second EC2 instance to process a failed instance's job
- An Elastic Load Balancer to send web traffic to healthy EC2 instances
- IAM user credentials on EC2 instances to grant permissions to modify an SQS queue
- An AutoScaling group to recover from EC2 instance failures

Explanation

"You are a consultant tasked with migrating an on-premise application architecture to AWS. During your design process you have to give consideration to current on-premise security and determine which security attributes you are responsible for on AWS. Which of the follow does AWS provide for you as part of the shared responsibility model? Choose the 2 correct answers:"

- Instance Security
- Physical Network Infrastructure
- Vitualization Infrastracture
- User access to the AWS envrionment

Explanation

Which statement is true about Amazon SQS? Choose the 2 correct answers:

- Amazon SQS (Simple Queue Service) guarantees delivery of AT LEAST 1 message but cannot guarantee it will not create duplicates
- Amazon SQS guarantees delivery of AT LEAST 1 message but cannot guarantee message order; but does attempt
- Amazon SQS guarantees delivery of AT LEAST 1 message and the message order which it is sent/received
- Amazon SQS (Simple Queue Service) guarantees delivery of AT LEAST 1 message and guarantees it will not create duplicates

Explanation

You create an SQS queue with the default settings for a new application your company is deploying. While new messages are added to the queue throughout the week, management has indicated that the application which retrieves the messages should only be run during your company's weekly Sunday evening maintenance window. It is quickly noticed on Monday morning that several messages were not processed the previous evening and the messages are no longer in the queue. What is a likely cause for this issue?

- Your application was using long polling, so some messages may not have been returned
- The EC2 instances the application runs on have permission to the queue, but not to some of the messages in the queue
- The messages surpassed the retention period for the queue
- The EC2 instances the application runs on are in a different subnet than the SQS queue

Explanation

You're building out an application on AWS that is running within a single region. However, you're designing with disaster recovery in mind. Your intention is to build the application so that if the current region it is running in becomes unavailable, you can failover to another region. Part of your application relies heavily on pre-built AMI's. In order to share those AMI's with the region you're using as a backup, what process would you take?

- Copy the AMI from the current region to another region, modify the auto scaling groups in the backup region to use the new AMI id in the backup region
- Modify the image permissions to be shared to the designated backup region
- Modify the image permissions to share the AMI with another account that set the default region to the backup region
- Nothing because all AMI's are available in any region as long as it is created within the same account

Explanation

What is the minimum size of an S3 object?

- 1Tb

- 0Byte
- 1GB
- 1Byte

Explanation

What is the maximum size of a general SSD EBS volume?

- 16TiB
- 4Gib
- 2TiB
- 16TB

Explanation

Data stored on EBS volumes are automatically and redundantly stored in multiple physical volumes in the same availability zone as part of the normal operations of the EBS service and at no additional charge.

- False
- True

Explanation

You are asked to review a plan that your company has made to create a new application that makes use of SQS, EC2, AutoScaling and CloudWatch. Which of the following action items will you advise your company not to implement?

- Utilize CloudWatch alarms to alert when the number of messages in the SQS queue grows too large
- Utilize an IAM role to grant EC2 instances permission to modify the SQS queue
- Utilize short polling with a wait time of 20 seconds to reduce the number of empty responses from the SQS queue
- Utilize AutoScaling to deploy new EC2 instances if the SQS queue grows too large

Explanation

After configuring a whole site CDN on CloudFront you receive the following error: This distribution is not configured to allow the HTTP request method that was used for this request. The distribution supports only cachable requests. What is the most likely cause of this?

- Allowed HTTP methods on that specific origin is only accepting GET, HEAD
- Allowed HTTP methods on that specific origin is only accepting GET, HEAD, OPTIONS
- The CloudFront distribution is configured to the wrong origin
- Allowed HTTP methods on that specific origin is only accepting GET, HEAD, OPTIONS, PUT, POST, PATCH, DELETE

Explanation

“You create an SQS queue and decide to test it out by creating a simple application which looks for messages in the queue. When a message is retrieved, the application is supposed to delete the message. You create three test messages in your SQS queue and discover that messages 1 and 3 are quickly deleted but message 2 remains in the queue. What is a possible cause for this behavior? Choose the 2 correct answers:”

- Message 2 uses JSON formatting.
- Your application is using short polling.
- The order that messages are received in is not guaranteed in SQS.
- You failed to set the correct permissions on message 2.

Explanation

Given the following region design, what is the bare minimum design for configuring high availability and why? Two availability zones within a region, az-a and az-b

- Create an ELB that serves traffic to one instances; two instances in each availability zone; enable cross-zone load balancing on the ELB; Because this applies high availability if a single availability zone is no longer available

- Create an ELB that serves traffic to two instances; one instance in each availability zone; enable cross-zone load balancing on the ELB; Create an Auto Scaling Group; Because this applies high availability if a single availability zone is no longer available
- Create an ELB that serves traffic to two instances; one instance in each availability zone; enable cross-zone load balancing on the ELB; Because this applies high availability if a single availability zone is no longer available
- Create an ELB that serves traffic to two instances; one instance in each availability zone; Because this applies high availability if a single availability zone is no longer available

Explanation

In AWS when a request is made, the AWS service decides whether a given request should be allowed or denied. The distinction between a request being denied or allowed by default and an explicit deny in a policy is important. Which of the following statements best describes this distinction?

- By default, a request is denied, but this can be overridden by an allow. In contrast, if a policy explicitly denies a request, that deny can't be overridden.
- By default, a request can either be allowed or denied depending on the request and each can be overwritten by the other. In contrast, if a policy explicitly denies a request, that deny can't be overridden.
- By default, a request is denied, but this can be overridden by an allow. In contrast, if a policy explicitly denies a request, it can be overridden but only by certain requests.
- By default, a request is allowed, but this can be overridden by a deny. In contrast, if a policy explicitly denies a request, that deny can't be overridden.

Explanation

While running an EC2 instance, you've been storing data on the root volume of the instance. However, in order to save money on the instance, you shut down the instance for the weekend. The following week, after starting the instance, you notice that all your work has been lost and is no longer available on the EC2 instance. What might be the cause of this?

- The EBS volume was not big enough to handle all of the processing data
- The EC2 instance was using instance store backed root volumes which are ephemeral and only live for the life of the instance
- The instance might have been compromised
- The EC2 instance was using EBS backed root volumes which are ephemeral and only live for the life of the instance

Explanation

After setting up you first VPC peering connection between you and one of your client's, the client requests that he now wants to be able to send traffic between instances in the peered VPCs using private IP addresses. What must you do to make this possible?

- Set up a second VPC peering connection
- Use IPSec tunneling
- Add his instance and your instance to the same placement group
- Add a route to a route table that's associated with your VPC

Explanation

You are designing a global application that takes advantage of multiple regions. As part of your application, the need to synchronize from one region to another is required to ensure your application is serving the same data when employing latency based Route 53 DNS records. In order to ensure this happens, you have determined that using the AWS CLI to sync files from the primary storage servers to S3 is the best method. How might you implement AWS CLI authentication against the S3 service?

- Assign API credentials to a standard user and configure the API keys in the AMI and copy the AMI to all the regions for deployment
- Create an EC2 IAM role and assign it to each EC2 instance that utilizes the AWS CLI to sync the data
- For each region, create an IAM EC2 instance role and assign it to the instance that needs to utilize the AWS CLI to sync the data from S3

- Create a cross account role and deploy the EC2 instances that need to sync data into separate AWS accounts

Explanation

Ten students have just been employed by your company for one week and it is your task to provide them with access to AWS, obviously through IAM. Your supervisor has come to you and said that he wants to be able to track these students as a group, for some bizarre reason, rather than individually. Because of this, he has requested for them to all have the same login ID but completely different passwords. Which of the following is the best way to achieve this? (Seems unlikely a supervisor will not be aware of the fact that you cannot have multiple IAM Users on the same account)

- It isn't possible to have the same login ID for multiple IAM users of the same accoun
- Use Multi Factor Authentication to verify each user and they will all be able to use the same login
- Create a separate login ID but give each IAM user the same alias so that each one can login with their alias
- Create various groups and add each user with the same login ID to different groups. The user can login with their own group ID

Explanation

Your supervisor asks you to create a decoupled application whose process includes dependencies on EC2 instances and servers located in your company's on-premises datacenter. Which of these are you least likely to recommend as part of that process?

- SQS polling from an EC2 instance deployed with an IAM role
- SQS polling from an EC2 instance using IAM user credentials
- An SWF workflow
- SQS polling from an on-premises server using IAM user credentials

Explanation

Which of the following is an invalid VPC Peering Configuration?

- You have a VPC peering connection between VPC A and VPC B . VPC A also has a VPN connection to a corporate network. You use VPC A to extend the peering relationship to exist between VPC B and the corporate network so that traffic from the corporate network can directly access VPC B by using the VPN connection to VPC A.
- You have a VPC peering connection between VPC A and VPC B, which are in the same AWS account, and do not have overlapping CIDR blocks.
- You have a central VPC (VPC A), and you have a VPC peering connection between VPC A and VPC B , and between VPC A and VPC C . The VPCs are in the same AWS account, and do not have overlapping CIDR blocks.
- You have peered three VPCs together in a full mesh configuration. The VPCs are in the same AWS account and do not have overlapping CIDR blocks

Explanation

You are consulting for a finance company that has specific backup and archiving policies. This company's RTO for all financial documents that have been created in the past 6 months is 1 hour. The second requirement is to configure a setup that allows for all documents that are 6 months or older to be sent automatically for archiving in a lower cost but highly durable archive environment. Given that the company is using the storage gateway, gateway-stored configuration, which of the following would be the best setup in order to reach the RTO and archiving objectives?

- Enable versioning on the S3 connected bucket to the gateway storage configuration
- Enable S3 lifecycle policy to immediately send all objects added to the bucket to Glacier
- Enable S3 versioning with a lifecycle policy to send objects older than 6 months to Amazon Glacier
- Enable S3 lifecycle policy to integrate into Amazon EBS for a good backup solution

Explanation

Company B provides an online image recognition service and utilizes SQS to decouple system components for scaleability. The SQS consumer's readers poll the image queue as often as possible to keep end-to-end throughput as high as possible. However, Company B is realizing that polling in tight loops is burning CPU cycles and increasing costs with empty responses. How can company B reduce the number of empty responses?

- Scale the component making the request using autoscaling based off the number of messages in the queue
- Enable short polling on the SQS message by setting the ReceiveMessageWaitTimeSeconds to a number = 0
- Enable long polling by setting the ReceiveMessageWaitTimeSeconds to a number > 0
- Enable short polling on the SQS queue by setting the ReceiveMessageWaitTimeSeconds to a number > 0

Explanation

Your application's usage peaks at 90% during the hours of 9am and 10am everyday. All other hours require only 10% of the peak resources. What is the best way to scale your application so that you're only paying for max resources during peak hours ?

- Run enough instances to handle peek capcity
- Auto Scaling by demand
- Proactive Cycle Scaling
- Proactive Event-based scaling

Explanation

As part of your application architecture requirements, the company you are working for has requested the ability to run analytics against all combined log files from the Elastic Load Balancer. Which services are used together in order to collect logs and process log file analysis in an AWS environment?

- Amazon S3 for storing ELB log files and Amazon EMR for processing the log files in analysis
- Amazon EC2 for storing and processing the log files
- Amazon DynamoDB to store the logs and EC2 for running custom log analysis scripts
- Amazon S3 for storing the ELB log files and EC2 for processing the log files in analysis

Explanation

Your supervisor calls you wanting to know why an SWF workflow you created has not made any progress in the last three weeks. What is the most likely explanation for the workflow's behaviour?

- The workflow has exceeded SWF's 14-day maximum workflow execution time
- SWF is awaiting human input from an activity task you assigned to your supervisor
- SWF does not support the tasks you created for on-premises servers, so the workflow has entered a paused state.
- The last task has exceeded SWF's 14-day maximum task execution time